

Aprendizaje supervisado

Aprendizaje supervisado

- Entregar entradas y salidas deseadas.
- Máquina produce algoritmo
- Predicciones para nuevas entradas.
- Ejemplos
 - Reconocer escritura a mano
 - Diagnóstico de tumores
 - Actividad fraudulenta en tarjetas de crédito


Clasificación y Regresión

- Clasificación
 - Predecir *etiquetas* de una lista predefinida de posibilidades.
 - Clasificación binaria
 - Clasificación multiclase
- Regresión
 - Predecir de una forma continua un número real

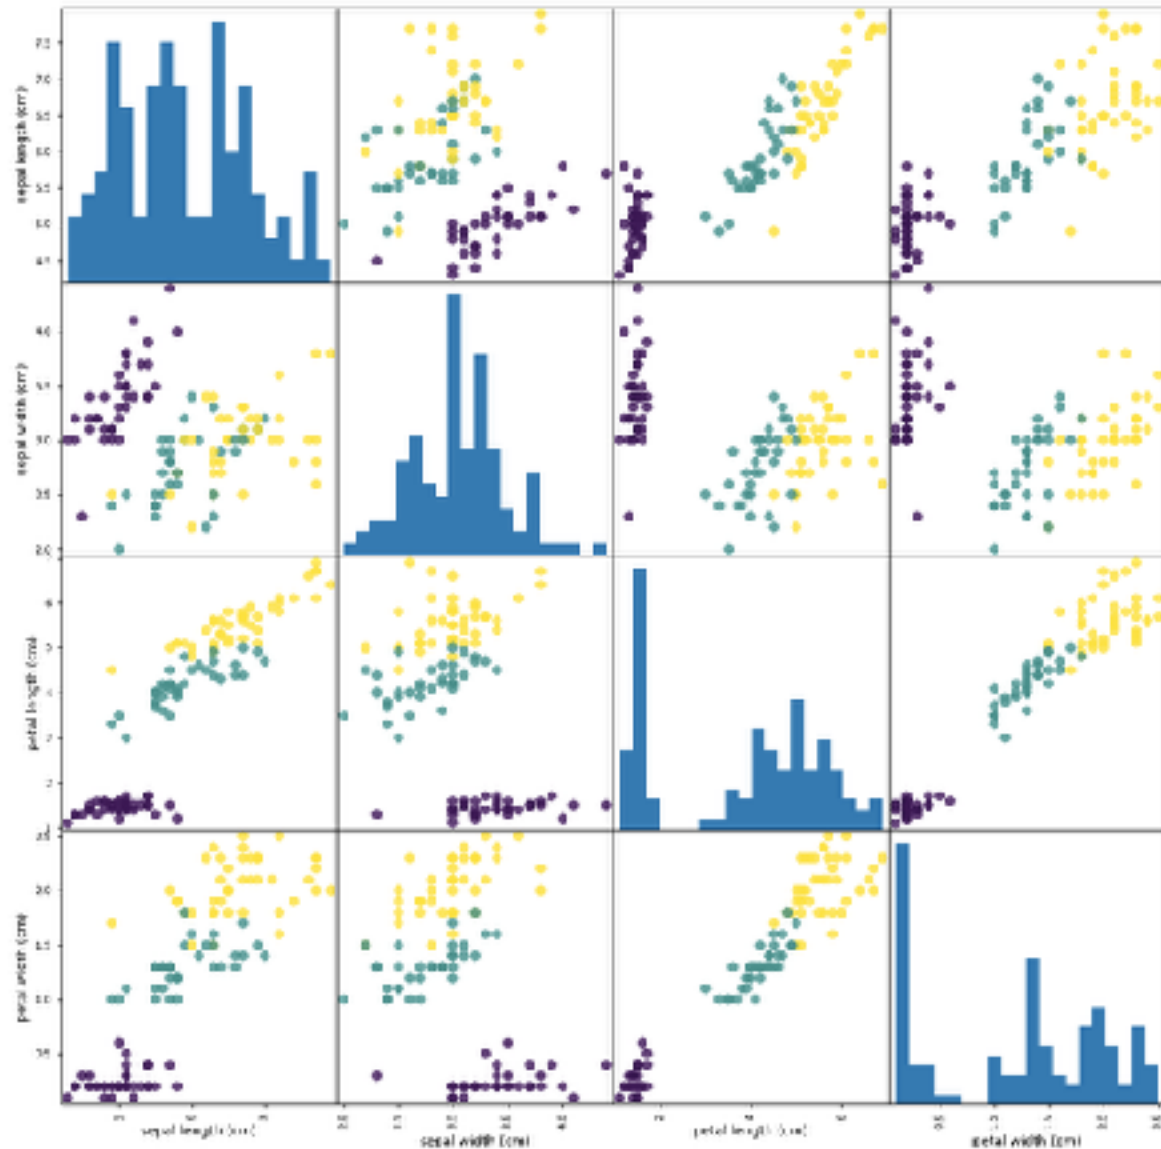
Petal



Sepal



Iris Data Set



K-Neighbors

```
from sklearn.datasets import load_iris  
iris_dataset = load_iris()
```

```
X_train, X_test, y_train, y_test = train_test_split(  
    iris_dataset['data'], iris_dataset['target'], random_state=0)
```

```
knn = KNeighborsClassifier(n_neighbors=1)
```

```
knn.fit(X_train, y_train)
```


```
print("Test set score: {:.2f}".format(knn.score(X_test, y_test)))
```

Test set score: 0.97

Reglas


Age	Number of cars owned	Owns house	Number of children	Marital status	Owns a dog	Bought a boat
66	1	yes	2	widowed	no	yes
52	2	yes	3	married	no	yes
22	0	no	0	married	yes	no
25	1	no	1	single	no	no
44	0	no	2	divorced	yes	no
39	1	yes	2	married	yes	no
26	1	no	2	single	no	no
40	3	yes	1	married	yes	no
53	2	yes	2	divorced	no	yes
64	2	yes	3	divorced	no	no
58	2	yes	2	married	yes	yes
33	1	no	1	single	no	no

Fitting




Clasificación

```
X, y = mglearn.datasets.make_forge()  
mglearn.discrete_scatter(X[:, 0], X[:, 1], y)  
plt.legend(["Class 0", "Class 1"], loc=4)  
plt.xlabel("First feature")  
plt.ylabel("Second feature")
```




Objetivos




```
X, y = mglearn.datasets.make_wave(n_samples=40)
plt.plot(X, y, 'o')
plt.ylim(-3, 3)
plt.xlabel("Feature")
plt.ylabel("Target")
```

k-Neighbors



`mglearn.plots.plot_knn_classification(n_neighbors=1)`

K-Neighbors



K- Neighbors

```
mglearn.plots.plot_knn_classification(n_neighbors=3)
```

```
from sklearn.model_selection import train_test_split
```

```
X, y = mglearn.datasets.make_forge()
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
clf = KNeighborsClassifier(n_neighbors=3)
```

```
clf.fit(X_train, y_train)
```


```
print("Test set predictions: {}".format(clf.predict(X_test)))
```

```
Test set predictions: [1 0 1 0 1 0 0]
```

```
print("Test set accuracy: {:.2f}".format(clf.score(X_test, y_test)))
```

```
Test set accuracy: 0.86
```

Análisis de k-neighbors



```
fig, axes = plt.subplots(1, 3, figsize=(10, 3))
```

```
for n_neighbors, ax in zip([1, 3, 9], axes):
```

```
# the fit method returns the object self, so we can instantiate
```

```
# and fit in one line
```

```
clf = KNeighborsClassifier(n_neighbors=n_neighbors).fit(X, y) mglearn.plots.plot_2d_separator(clf, X, fill=True, eps=0.5, ax=ax, alpha=.4)
```

```
mglearn.discrete_scatter(X[:, 0], X[:, 1], y, ax=ax)
```


```
ax.set_title("{} neighbor(s)".format(n_neighbors))
```

```
ax.set_xlabel("feature 0")
```


```
ax.set_ylabel("feature 1")
```

```
axes[0].legend(loc=3)
```

Precisión entrenamiento y prueba




k-vecinos (k=1) Regresión




`mglearn.plots.plot_knn_regression(n_neighbors=3)`

k-vecinos (k=3) Regresión



`mglearn.plots.plot_knn_regression(n_neighbors=3)`

Comparación regresiones




Modelos lineales

$$\hat{y} = w[0] * x[0] + w[1] * x[1] + \dots + w[p] * x[p] + b$$

Ejemplo de modelo lineal

$$\hat{y} = w[0] * x[0] + b$$

w[0]: 0.393906 b: -0.031804



Mínimos cuadrados

```
from sklearn.linear_model import LinearRegression
X, y = mglearn.datasets.make_wave(n_samples=60)
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)

lr = LinearRegression().fit(X_train, y_train)
print("lr.coef_: {}".format(lr.coef_))

print("lr.intercept_: {}".format(lr.intercept_))

lr.coef_: [ 0.394]
lr.intercept_: -0.031804343026759746

print("Training set score: {:.2f}".format(lr.score(X_train, y_train)))
print("Test set score: {:.2f}".format(lr.score(X_test, y_test)))

Training set score: 0.67
Test set score: 0.66
```

Regresión ridge

- Modelo lineal de regresión
- Coeficientes tan pequeños como sea posible
- Regularización para evitar *overfitting*

```
from sklearn.linear_model import Ridge
```

```
ridge = Ridge().fit(X_train, y_train)
```

```
print("Training set score: {:.2f}".format(ridge.score(X_train, y_train)))
```

```
print("Test set score: {:.2f}".format(ridge.score(X_test, y_test)))
```

Training set score: 0.89

Test set score: 0.75

```
ridge10 = Ridge(alpha=10).fit(X_train, y_train)
```

```
print("Training set score: {:.2f}".format(ridge10.score(X_train, y_train)))
```

```
print("Test set score: {:.2f}".format(ridge10.score(X_test, y_test)))
```

Training set score: 0.79

Test set score: 0.64

```
ridge01 = Ridge(alpha=0.1).fit(X_train, y_train)
```


```
print("Training set score: {:.2f}".format(ridge01.score(X_train, y_train)))
```

```
print("Test set score: {:.2f}".format(ridge01.score(X_test, y_test)))
```


Training set score: 0.93

Test set score: 0.77

Comparación regresión ridge



Curvas de aprendizaje



Lasso

- Regularización L1 - algunos coeficientes = 0

```
from sklearn.linear_model import Lasso
```

```
lasso = Lasso().fit(X_train, y_train)
```

```
print("Training set score: {:.2f}".format(lasso.score(X_train, y_train))) print("Test set score: {:.2f}".format(lasso.score(X_test, y_test)))
```


```
print("Number of features used: {}".format(np.sum(lasso.coef_ != 0)))
```

Training set score: 0.29

Test set score: 0.21

Number of features used: 4

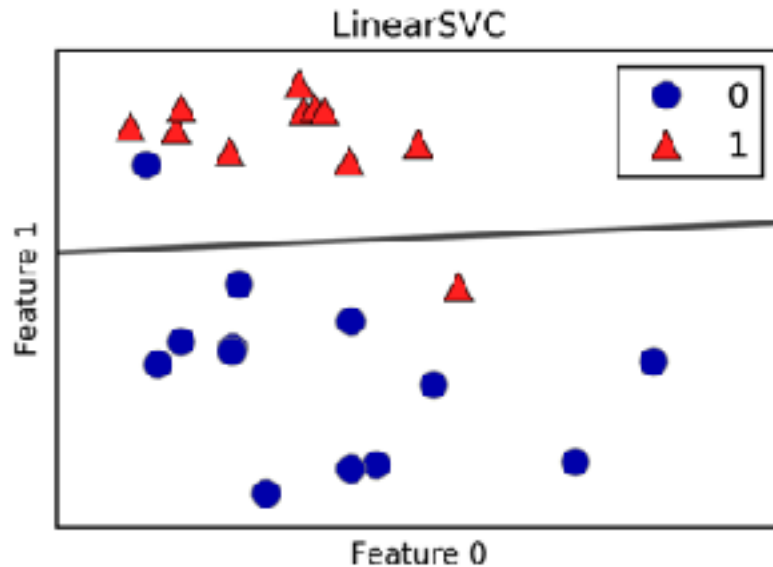
Lasso comparado con Ridge



Clasificación lineal

$$\hat{y} = w[0] * x[0] + w[1] * x[1] + \dots + w[p] * x[p] + b > 0$$


Vector de soporte (SVC) y Regresión logística



```
X, y = mglearn.datasets.make_forge()
fig, axes = plt.subplots(1, 2, figsize=(10, 3))
for model, ax in zip([LinearSVC(), LogisticRegression()], axes):
    clf = model.fit(X, y)
    mglearn.plots.plot_2d_separator(clf, X, fill=False, eps=0.5,
                                   ax=ax, alpha=.7)
    mglearn.discrete_scatter(X[:, 0], X[:, 1], y, ax=ax)
ax.set_title("{}".format(clf.__class__.__name__))
ax.set_xlabel("Feature 0")
ax.set_ylabel("Feature 1")

axes[0].legend()
```

Sensibilidad a C - SVC



`mglearn.plots.plot_linear_svc_regularization()`

Sensibilidad a C - Regresión logística


```
from sklearn.datasets import load_breast_cancer cancer =  
load_breast_cancer()  
X_train, X_test, y_train, y_test = train_test_split(  
cancer.data, cancer.target, stratify=cancer.target, random_state=42)  
logreg = LogisticRegression().fit(X_train, y_train)  
print("Training set score: {:.3f}".format(logreg.score(X_train,  
y_train))) print("Test set score: {:.3f}".format(logreg.score(X_test,  
y_test)))
```

Training set score: 0.953

Test set score: 0.958

```
logreg100 = LogisticRegression(C=100).fit(X_train, y_train)  
print("Training set score: {:.3f}".format(logreg100.score(X_train,  
y_train))) print("Test set score: {:.  
3f}".format(logreg100.score(X_test, y_test)))
```

```
plt.plot(logreg.coef_.T, 'o', label="C=1")  
plt.plot(logreg100.coef_.T, '^', label="C=100")  
plt.plot(logreg001.coef_.T, 'v', label="C=0.001")  
plt.xticks(range(cancer.data.shape[1]), cancer.feature_names,  
rotation=90)  
plt.hlines(0, 0, cancer.data.shape[1])  
plt.ylim(-5, 5)  
plt.xlabel("Coefficient index")  
plt.ylabel("Coefficient magnitude")  
plt.legend()
```



Coeficientes con penalidad

```
for C, marker in zip([0.001, 1, 100], ['o', '^', 'v']):  
    lr_l1 = LogisticRegression(C=C, penalty="l1").fit(X_train,  
    y_train) print("Training accuracy of l1 logreg with C={:.3f}: {:.  
    2f}".format(
```

```
C, lr_l1.score(X_train, y_train)))  
print("Test accuracy of l1 logreg with C={:.3f}: {:.2f}".format(
```

```
    C, lr_l1.score(X_test, y_test)))  
    plt.plot(lr_l1.coef_.T, marker, label="C={:.3f}".format(C))  
    plt.xticks(range(cancer.data.shape[1]), cancer.feature_names,  
    rotation=90)
```

```
plt.hlines(0, 0, cancer.data.shape[1])
```

```
plt.xlabel("Coefficient index")
```

```
plt.ylabel("Coefficient magnitude")
```

```
plt.ylim(-5, 5)
```

```
plt.legend(loc=3)
```

Training accuracy of l1 logreg with C=0.001: 0.91


Test accuracy of l1 logreg with C=0.001: 0.92

Training accuracy of l1 logreg with C=1.000: 0.96


Test accuracy of l1 logreg with C=1.000: 0.96

Training accuracy of l1 logreg with C=100.000: 0.99


Test accuracy of l1 logreg with C=100.000: 0.98




Clasificadores lineales - multiclase



Predicción multiclase - SVC



Decisiones multiclase



Conclusiones

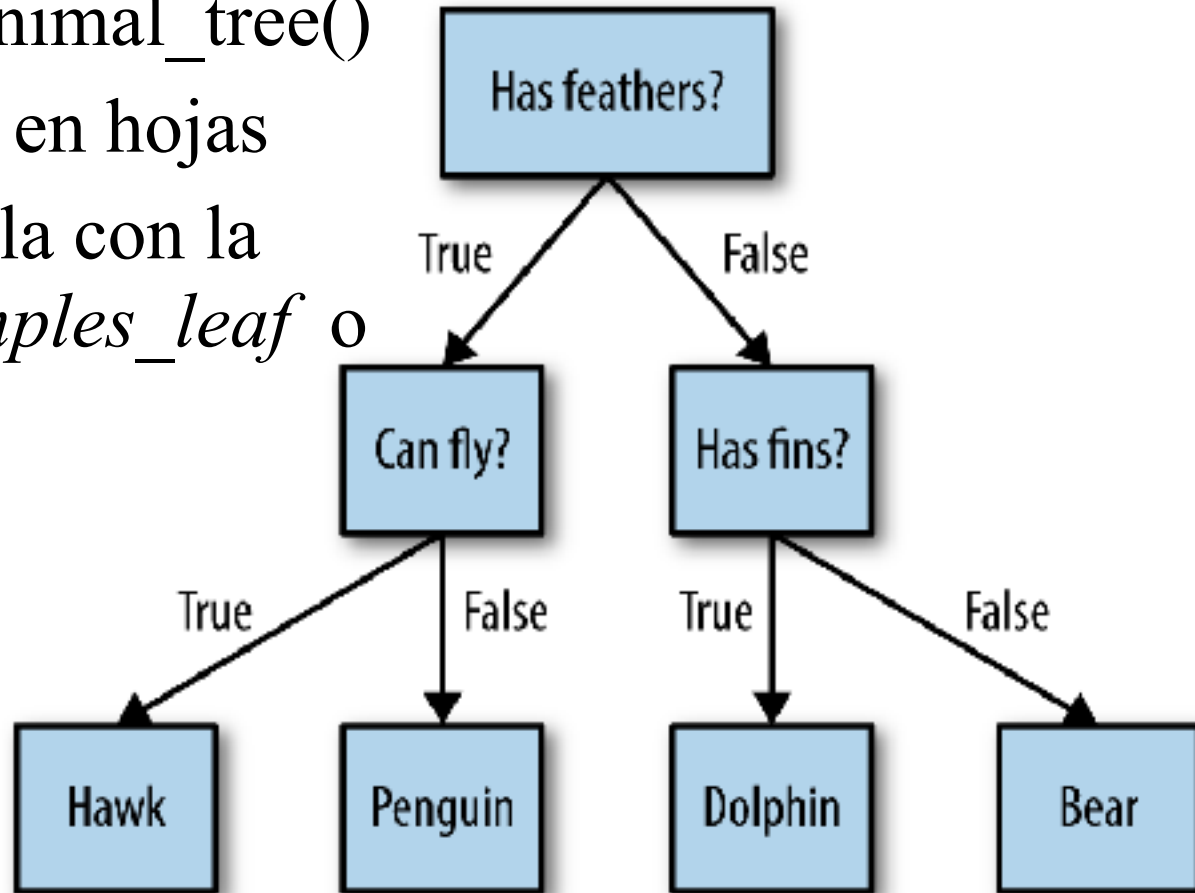
- Principal parámetro es la regularización
- Alto valor de α o bajo valor de C -> modelos más simples
- Si pocas características son importantes use regularización L1.
- Pocas características son más simples de explicar.
- Modelos lineales son rápidos para entrenar y predecir.
- Para bases datos muy grandes solver='sag'
-

Naive Bayes


- Rápidos en entrenamiento
- Desempeño en generalización inferior a modelos lineales
- Útiles para datasets grandes de múltiples dimensiones
- Usan un parámetro α
- Tipos de Naive Bayes
 - GaussianNB - datos continuos
 - BernoulliNB - datos binarios
 - MultinomialNB - cuentas enteros

Arboles de decisión

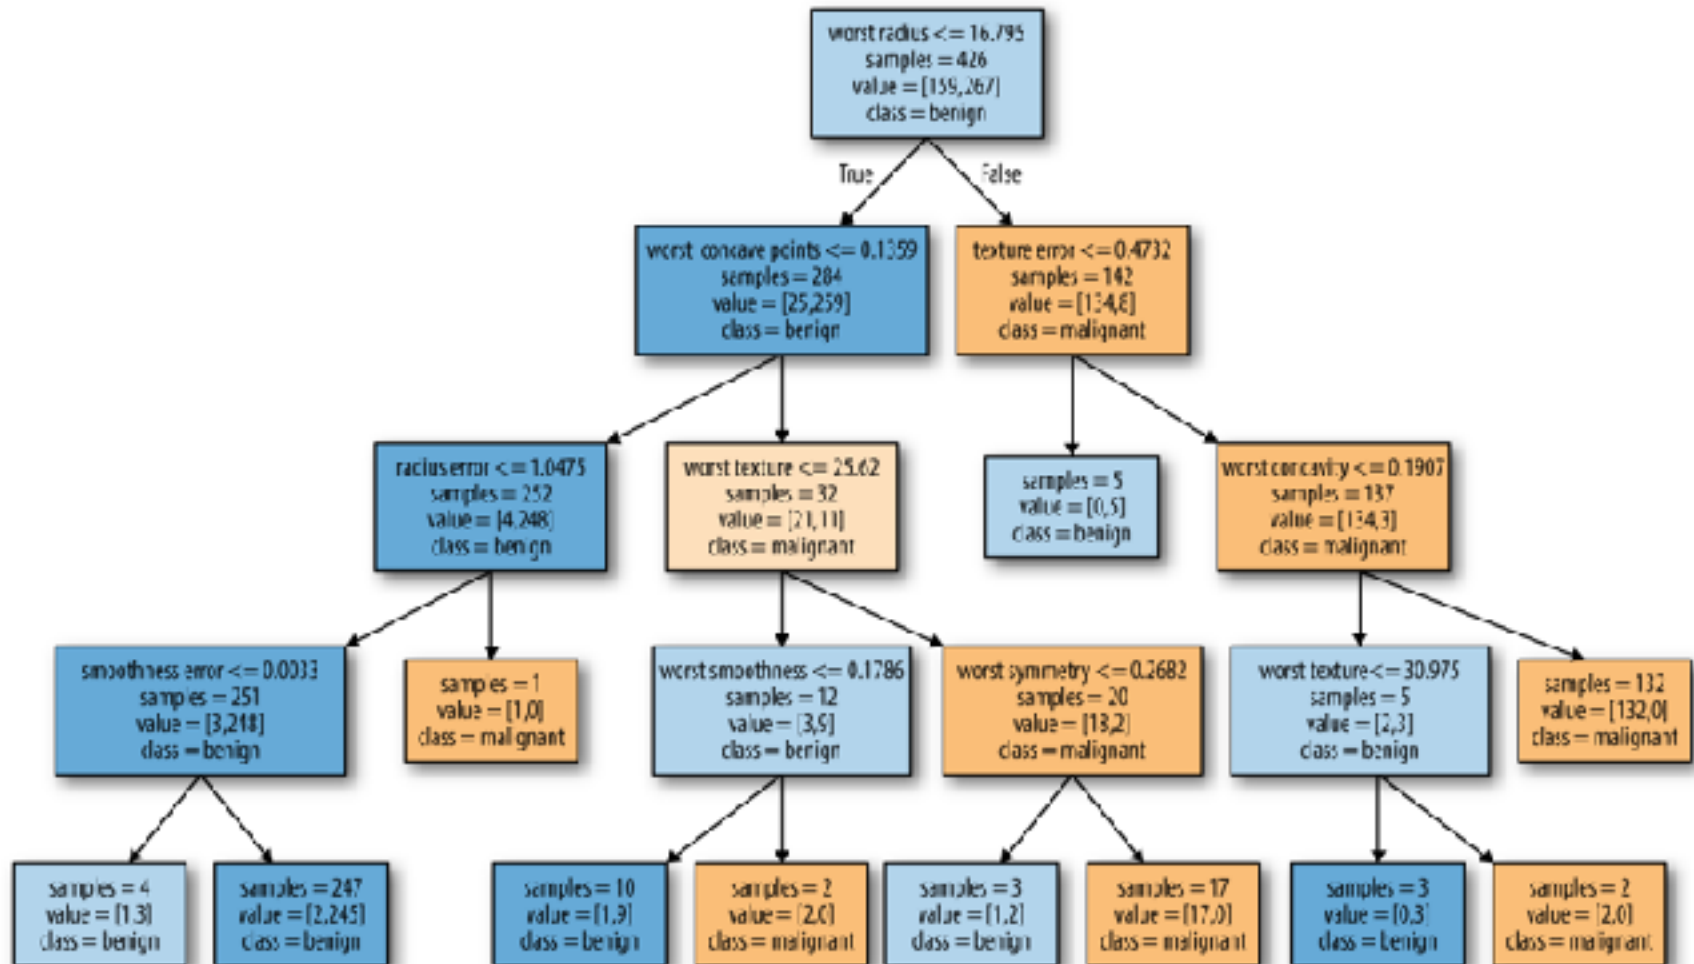
- Estructura if/else
- `mglearn.plots.plot_animal_tree()`
- Preguntas o terminal en hojas
- Overfitting se controla con la *max_depth*, *min_samples_leaf* o *max_leaf_nodes*
- Se pueden analizar
- Random forest




Arboles de decisión




Análisis del árbol de decisiones









Importancia de las características



Comparación




Random forest



Gradient boosted regression

- Construye árboles serialmente
- Cada árbol trata de corregir errores de los anteriores
- Poca profundidad - 1 a 5 lo cual requiere menos memoria y mejora desempeño.
- Árboles agregados iterativamente mejoran desempeño.
- Se controla *learning_rate* que tan fuerte se corrigen los errores anteriores. Un valor alto hace el modelo más complejo
- Desempeño es superior al random forest
- Toma tiempo para entrenar
- Sensible a los parámetros
 - Número de árboles *n_estimators*
 - *learning_rate*

Features



Support Vector Machines - SVM

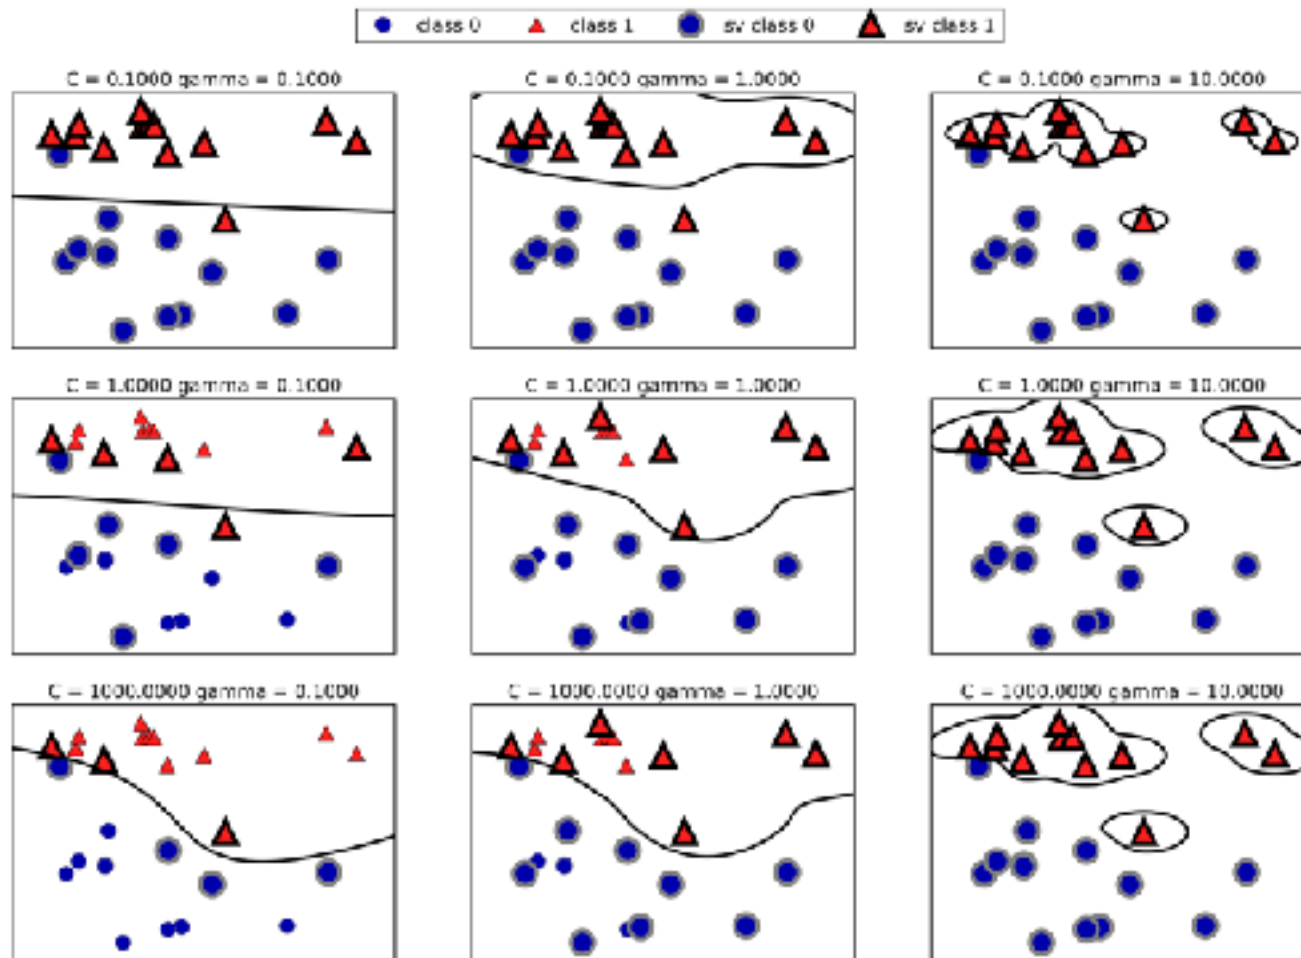
- Los datos del borde de decisión son más importantes. Son el *support vector*
- Para hacer una predicción, la distancia a cada vector del punto es medida.
- $k_{\text{rbf}}(x_1, x_2) = \exp(-\gamma \|x_1 - x_2\|^2)$
- La importancia de los vectores es el **aprendizaje**
- y ancho de la gaussiana

□

C y Gamma (γ) en SVM



- *Gamma* limita el ancho del Gaussiano
- *Gamma* pequeño significa un radio grande y muchos puntos cercanos son considerados.
- *Gamma* bajo produce un modelo menos complejo
- C es el parámetro de regularización. Limita la importancia de cada punto.
- C más alto hace que el borde de decisión los clasifique correctamente.

Support Vector Machines (SVM)



Redes neuronales

Inputs




$$h[0] = \tanh(w[0, 0] * x[0] + w[1, 0] * x[1] + w[2, 0] * x[2] + w[3, 0] * x[3])$$

$$h[1] = \tanh(w[0, 0] * x[0] + w[1, 0] * x[1] + w[2, 0] * x[2] + w[3, 0] * x[3])$$

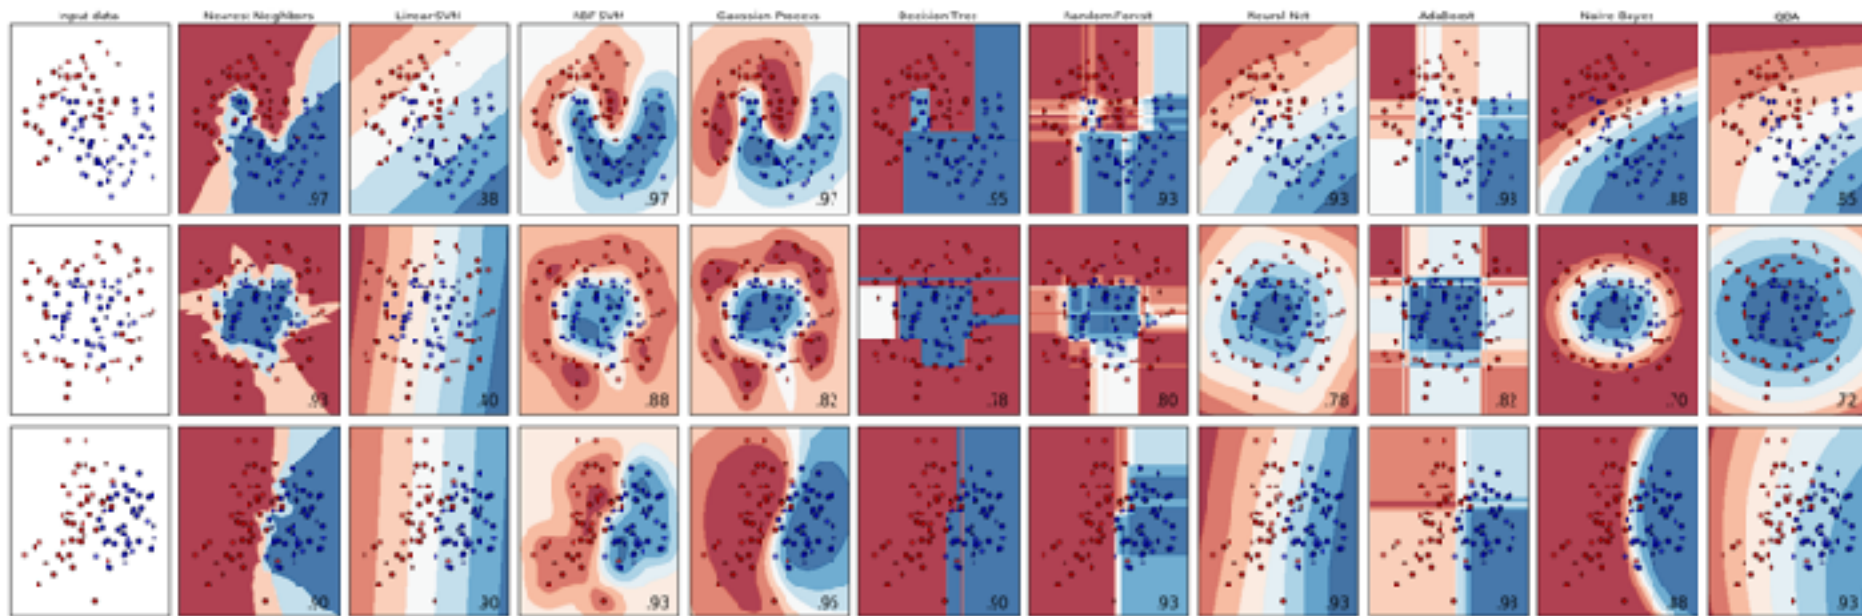
$$h[2] = \tanh(w[0, 0] * x[0] + w[1, 0] * x[1] + w[2, 0] * x[2] + w[3, 0] * x[3])$$

$$\hat{y} = v[0] * h[0] + v[1] * h[1] + v[2] * h[2]$$

Redes neuronales



Comparación de modelos



Comparación

- Nearest neighbors
- Linear models
- Naive Bayes
- Decision trees
- Random forests
- Gradient boosted decision trees
- Support vector machines
- Neural networks